

---

# ML Optimizers: 10 Essential Types

[Abdelrahman Hamdy Mustafa Ali El-Hamshary | 202002570](#)

This list details ten types of optimizers used in Machine Learning, each with a brief functional explanation:

1. **Gradient Descent (GD):** A core optimization algorithm that adjusts model parameters repeatedly to minimize the loss function by moving along the negative gradient direction<sup>2</sup>.
  2. **Stochastic Gradient Descent (SGD):** An iteration of GD that speeds up convergence by updating parameters based on a randomly selected subset (mini-batch) of the training data<sup>3</sup>.
  3. **Mini-batch Gradient Descent:** This method balances speed and stability by using small batches of data for parameter updates, combining elements of both GD and SGD<sup>4</sup>.
  4. **Adagrad:** Adapts the learning rate individually for each parameter based on its past gradients, making it particularly useful for sparse data applications<sup>5</sup>.
  5. **RMSprop:** An adaptive learning rate method that divides the gradient by a running average of the magnitudes of recent gradients. It's well-suited for non-stationary objectives<sup>6</sup>.
  6. **Adam (Adaptive Moment Estimation):** A highly effective optimizer that integrates both adaptive learning rates and momentum for fast convergence and efficient training<sup>7</sup>.
  7. **AdaDelta:** Extends Adagrad by restricting the accumulation of past squared gradients, thereby preventing the learning rate from decaying too aggressively over time<sup>8</sup>.
  8. **Nadam:** A variant of the Adam optimizer that incorporates Nesterov momentum, which can lead to potentially faster model convergence<sup>9</sup>.
  9. **Momentum:** Accelerates the SGD process by adding a portion of the prior update to the current parameter update, significantly boosting convergence speed<sup>10</sup>.
  10. **Newton's Method:** Finds optimal parameters by approximating the curvature of the loss surface using second-order derivatives (the Hessian matrix). This approach is generally more computationally costly<sup>11</sup>.
-